

# week4

April 12, 2024

## 0.1 Week 4 : NLP Disaster Tweets Kaggle Mini-Project

Student, University of Colorado Boulder

## 0.2 Step 1: Brief description of the problem and data

We are provided with tweets which are classified as 1 if they are about real disasters and 0 if they are not. We will use these as training data. Intent of the project is to build a machine learning model that predicts which tweets are about real disasters and which one's aren't. We will develop this model and test them on the test data.

The details of the project is available in Kaggle competition <https://www.kaggle.com/c/nlp-getting-started/overview>

```
[1]: # Imports

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import nltk
from nltk.stem import WordNetLemmatizer

from wordcloud import WordCloud, STOPWORDS, ImageColorGenerator

from copy import deepcopy

from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
from tensorflow.keras.losses import BinaryCrossentropy
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.metrics import AUC
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.models import load_model
```

```
/Users/kkuruva/Desktop/masters/UniversityOfColoradoBoulder/DTSA 5511
Introduction to Deep Learning/week4/venv/lib/python3.9/site-
packages/urllib3/__init__.py:35: NotOpenSSLWarning: urllib3 v2 only supports
OpenSSL 1.1.1+, currently the 'ssl' module is compiled with 'LibreSSL 2.8.3'.
See: https://github.com/urllib3/urllib3/issues/3020
warnings.warn(
```

**Data Extraction** User has to register on kaggle to get access to the dataset. This data is expected to be stored in data/nlp-getting-started folder

```
[2]: data_path = 'data/nlp-getting-started/'
df_train = pd.read_csv(data_path+'train.csv')
df_test = pd.read_csv(data_path+'test.csv')

df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7613 entries, 0 to 7612
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0    id         7613 non-null   int64
 1   keyword    7552 non-null   object
 2   location    5080 non-null   object
 3    text       7613 non-null   object
 4   target     7613 non-null   int64
dtypes: int64(2), object(3)
memory usage: 297.5+ KB
```

There are total of 7613 rows with 5 columns. 2 columns are of integer type and 3 are of type object (string). The id field is unique with values going from 0 to 7612.

```
[3]: df_train.head(10)
```

```
[3]:
```

	id	keyword	location	text \	target
0	1	NaN	NaN	Our Deeds are the Reason of this #earthquake M...	1
1	4	NaN	NaN	Forest fire near La Ronge Sask. Canada	1
2	5	NaN	NaN	All residents asked to 'shelter in place' are ...	1
3	6	NaN	NaN	13,000 people receive #wildfires evacuation or...	1
4	7	NaN	NaN	Just got sent this photo from Ruby #Alaska as ...	1
5	8	NaN	NaN	#RockyFire Update => California Hwy. 20 closed...	1
6	10	NaN	NaN	#flood #disaster Heavy rain causes flash flood...	1
7	13	NaN	NaN	I'm on top of the hill and I can see a fire in...	1
8	14	NaN	NaN	There's an emergency evacuation happening now ...	1
9	15	NaN	NaN	I'm afraid that the tornado is coming to our a...	1

```
8      1
9      1
```

```
[4]: df_train.describe()
```

```
[4]:
```

	id	target
count	7613.000000	7613.000000
mean	5441.934848	0.42966
std	3137.116090	0.49506
min	1.000000	0.00000
25%	2734.000000	0.00000
50%	5408.000000	0.00000
75%	8146.000000	1.00000
max	10873.000000	1.00000

```
[5]: df_test.info
```

```
[5]: <bound method DataFrame.info of          id keyword location \
0         0      NaN      NaN
1         2      NaN      NaN
2         3      NaN      NaN
3         9      NaN      NaN
4        11      NaN      NaN
...
3258  10861      NaN      NaN
3259  10865      NaN      NaN
3260  10868      NaN      NaN
3261  10874      NaN      NaN
3262  10875      NaN      NaN

                                text
0                Just happened a terrible car crash
1    Heard about #earthquake is different cities, s...
2    there is a forest fire at spot pond, geese are...
3                Apocalypse lighting. #Spokane #wildfires
4        Typhoon Soudelor kills 28 in China and Taiwan
...
3258  EARTHQUAKE SAFETY LOS ANGELES Û SAFETY FASTE...
3259  Storm in RI worse than last hurricane. My city...
3260  Green Line derailment in Chicago http://t.co/U...
3261  MEG issues Hazardous Weather Outlook (HWO) htt...
3262  #CityofCalgary has activated its Municipal Eme...

[3263 rows x 4 columns]>
```

```
[6]: df_test.describe()
```

```
[6]:
```

	id
count	3263.000000
mean	5427.152927
std	3146.427221
min	0.000000
25%	2683.000000
50%	5500.000000
75%	8176.000000
max	10875.000000

We have 3263 rows of data in the test dataframe and the structure is similar to the training data without the target column

### 0.3 Step 2: Exploratory Data Analysis (EDA) — Inspect, Visualize and Clean the Data

Let us look at distribution of disaster vs not disaster tweets in the training data

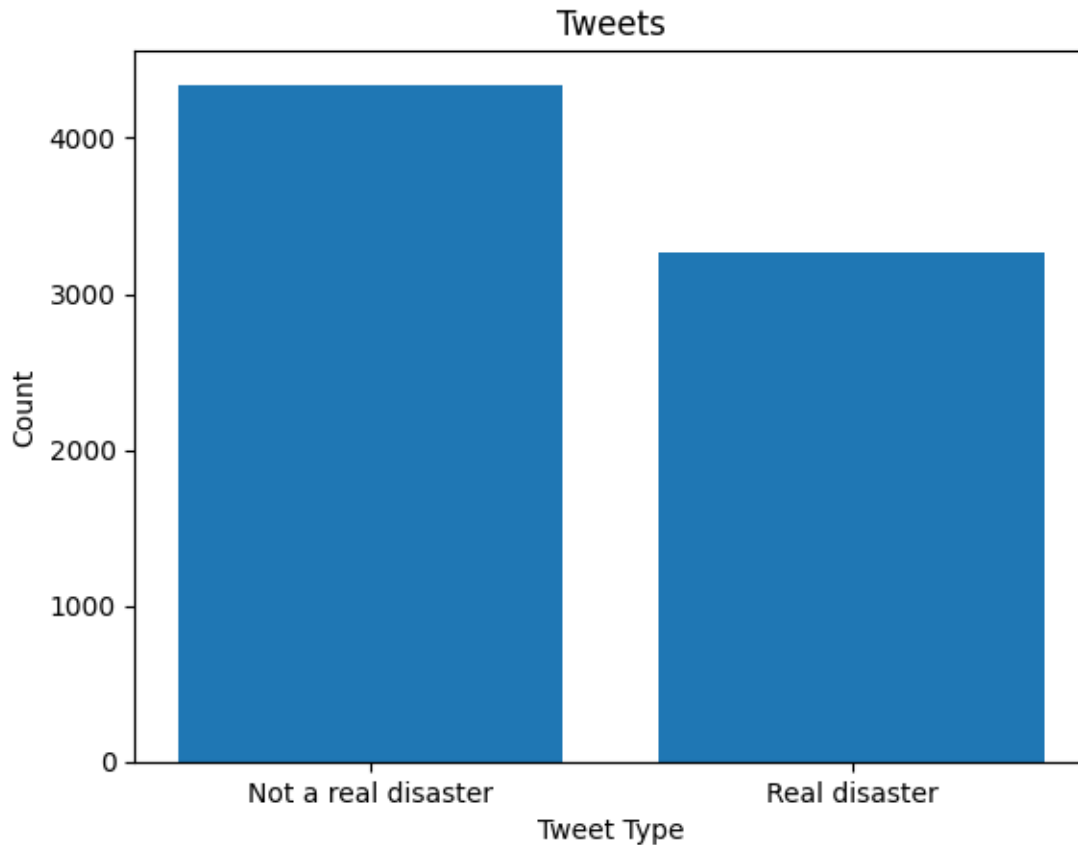
```
[7]: df_train['target'].plot.hist()

x = ["Not a real disaster", "Real disaster"]
y = [df_train['target'].value_counts()[0.0], df_train['target'].
     ↪value_counts()[1.0]]

plt.bar(x, y)
plt.title("Tweets")

plt.xlabel("Tweet Type")
plt.ylabel("Count")

plt.show()
```



Let us look at the size of tweets

```
[8]: def histogram_of_word_counts(strings, df_type):
    counts = {}
    for string in strings:
        words = string.split(' ')
        length = len(words)
        if length not in counts:
            counts[length] = 0
        counts[length] += 1
        #if length > 50:
        #    print(f'{string}: {words}')


    word_counts = list(counts.keys())

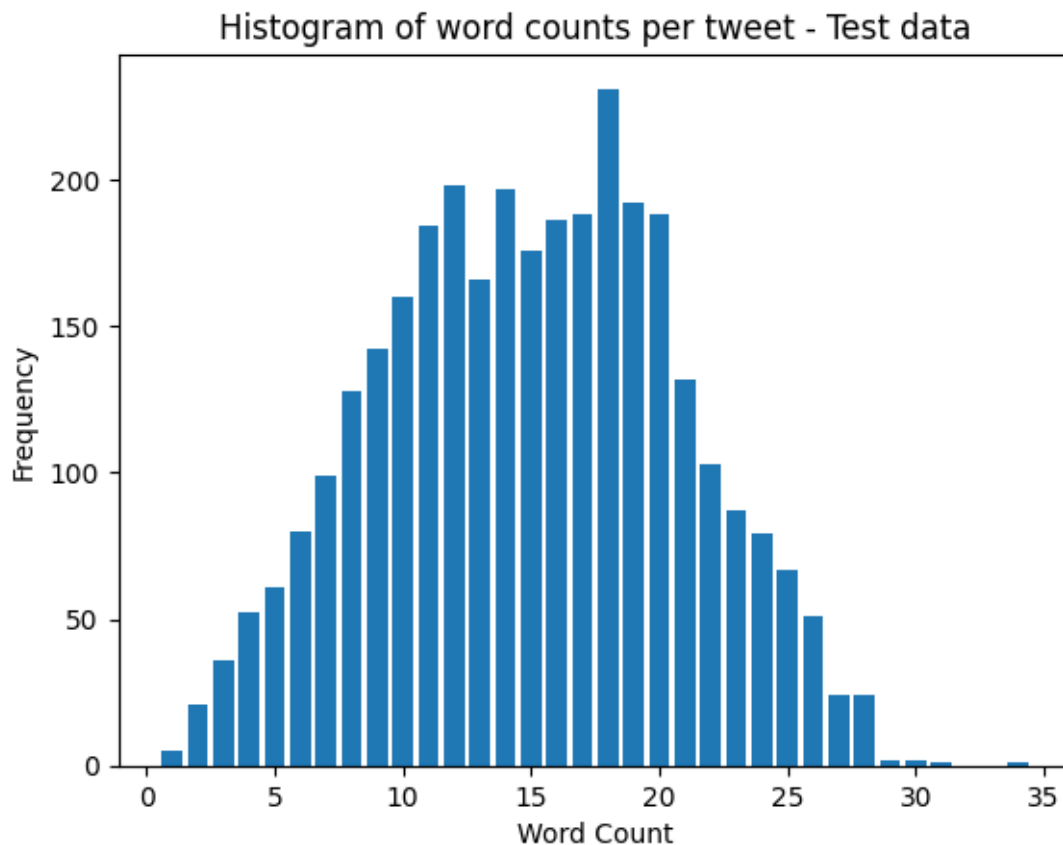
    counts = list(counts.values())

    # Create the histogram.
    plt.bar(word_counts, counts)
    plt.xlabel("Word Count")
```

```
plt.ylabel("Frequency")
plt.title(f"Histogram of word counts per tweet - {df_type} data")
plt.show()
```

```
histogram_of_word_counts(df_train['text'], 'Training')
histogram_of_word_counts(df_test['text'], 'Test')
```





Let us look at some frequent words in the tweets

```
[9]: stopwords = set(STOPWORDS)
stopwords.update(['u', 't', 'https', 'co', 'û_', 'û'])
text = " ".join(review for review in df_train.text)
# Generate a word cloud image
wordcloud = WordCloud(stopwords=stopwords, background_color="white").
    ↳generate(text)

# Display the generated image:
# the matplotlib way:
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")
plt.show()
```





```
[26]: !wget http://downloads.cs.stanford.edu/nlp/data/glove.6B.zip

--2024-04-10 10:18:21--  http://downloads.cs.stanford.edu/nlp/data/glove.6B.zip
Resolving downloads.cs.stanford.edu (downloads.cs.stanford.edu)... 171.64.64.22
Connecting to downloads.cs.stanford.edu
(downloads.cs.stanford.edu)|171.64.64.22|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 862182613 (822M) [application/zip]
Saving to: 'glove.6B.zip.1'

glove.6B.zip.1      100%[=====>] 822.24M   834KB/s   in 14m 26s

2024-04-10 10:32:46 (972 KB/s) - 'glove.6B.zip.1' saved [862182613/862182613]
```

3b. Extract the zip file

```
[18]: !unzip glove.6B.zip
```

```
Archive:  glove.6B.zip
  inflating: glove.6B.50d.txt
  inflating: glove.6B.100d.txt
  inflating: glove.6B.200d.txt
  inflating: glove.6B.300d.txt
```

3c. Form the GloVe word mappings. Let us use the 100 dimension file.

```
[11]: def glove_word_mappings(filename):
        word_mapping = dict()
        with open(filename, 'r') as f:
            for line in f.readlines():
                try:
                    words = line.split(' ')
                    word_mapping[words[0]] = np.array(words[1:], dtype=float)
                except Exception as ex:
                    print(f'failed to use {words[0]}: {ex}')
                    break
        return word_mapping

glove_words = glove_word_mappings('glove.6B.100d.txt')
print(f'count of glove words {len(glove_words)}')
```

```
count of glove words 400000
```

3d. Download wordnet using nltk. We will be using this for the tokenizer

```
[12]: nltk.download('wordnet')

regex_tokenizer = nltk.RegexpTokenizer(r'\w+')
```

```
[nltk_data] Downloading package wordnet to
[nltk_data]    /Users/kkuruvad/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
```

3e. Function to lemmatize the tweet/a given text. We will only use the lemmatized token if it is present in GloVe

```
[13]: lemmatizer = WordNetLemmatizer()

def tweet_to_tokens(tweet):
    tokens = regexp_tokenizer.tokenize(tweet)
    tokens = [t.lower() for t in tokens]
    lemmatized_tokens = [lemmatizer.lemmatize(t) for t in tokens]
    return [t for t in lemmatized_tokens if t in glove_words]
```

3f. Function to convert the tokens into vectors

```
[14]: def tokens_to_vectors(tokens):
    vectors = list()
    for token in tokens:
        if token not in glove_words:
            continue
        vectors.append(glove_words[token])
    return np.array(vectors, dtype=float)
```

3g. Split input data into training (70%) and validation (30%)

```
[15]: m_df_train = df_train.sample(frac=1, random_state=142)
m_df_train.reset_index(drop=True, inplace=True)

train_validate_split = int(len(m_df_train)*0.7)

train, validate = m_df_train[:train_validate_split],
↳ m_df_train[train_validate_split:]
len(train), len(validate)
```

```
[15]: (5329, 2284)
```

3h. Convert the training, validation and test dataframes into tokens

```
[16]: def dataframe_to_X_y(df):
    y = None
    if 'target' in df:
        y = df['target'].to_numpy().astype(int)

    vectors = list()
    for tweet in df['text']:
        tokens = tweet_to_tokens(tweet)
        v = tokens_to_vectors(tokens)
        if v.shape[0] == 0:
```

```

        vectors.append(np.zeros(shape=(1,100)))
    else:
        vectors.append(v)
    return vectors, y

df_test.info()
X_train, y_train = dataframe_to_X_y(train)
X_val, y_val = dataframe_to_X_y(validate)
X_test, _ = dataframe_to_X_y(df_test)

print(f'len(X_train): {len(X_train)}, len(y_train): {len(y_train)}, len(X_val): {len(X_val)}, len(y_val): {len(y_val)}, len(X_test): {len(X_test)}')

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3263 entries, 0 to 3262
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   id      3263 non-null    int64
 1   text    3263 non-null    object
dtypes: int64(1), object(1)
memory usage: 51.1+ KB
len(X_train): 5329, len(y_train): 5329, len(X_val): 2284, len(y_val): 2284,
len(X_test): 3263

```

3i. Check for the maximum vectors size of the tweet in all the 3 dataframes (training, validation and test) . We will use this later in padding

```

[17]: sequence_lengths = list()
      for i in range(len(X_train)):
          sequence_lengths.append(len(X_train[i]))
      for i in range(len(X_val)):
          sequence_lengths.append(len(X_val[i]))
      for i in range(len(X_test)):
          sequence_lengths.append(len(X_test[i]))


      plt.hist(sequence_lengths)
      pd.Series(sequence_lengths).describe()

```

```

[17]: count    10876.000000
      mean      15.636539
      std       5.927161
      min       1.000000
      25%      12.000000
      50%      16.000000
      75%      20.000000
      max      34.000000
      dtype: float64

```



3j. Use the maximum size of vectors to determine padding. We will round the value and consider 35

```
[18]: def pad_X(X, length, dimension):
        X_copy = deepcopy(X)
        for i,x in enumerate(X):
            l = x.shape[0]
            padding = np.zeros(shape=(length-x.shape[0], dimension))
            X_copy[i] = np.concatenate([x, padding])
        return np.array(X_copy).astype(float)

#X_train.shape
X_train_pad = pad_X(X_train, 35, 100)
print(f'training: {X_train_pad.shape}')

X_val_pad = pad_X(X_val, 35, 100)
print(f'validation: {X_val_pad.shape}')

X_test_pad = pad_X(X_test, 35, 100)
print(f'test: {X_test_pad.shape}')
```

```
training: (5329, 35, 100)
validation: (2284, 35, 100)
```

test: (3263, 35, 100)

**Models** Now that our data is converted, we can start with tensorflow models

After padding our X has shape of 35, 100. So lets declare input layer of this shape LSTM of 64 units is good enough in most cases. So lets add that. Let us also add a drop out layer with 0.2.

Let us repeat the same layers again and finally flatten and apply sigmoid activation since we want the output to be 0 or 1

```
[19]: model = Sequential([])

model.add(layers.Input(shape=(35,100)))
model.add(layers.LSTM(64, return_sequences=True))
model.add(layers.Dropout(0.2))
model.add(layers.LSTM(64, return_sequences=True))
model.add(layers.Dropout(0.2))
model.add(layers.Flatten())
model.add(layers.Dense(1, activation='sigmoid'))

cp = ModelCheckpoint('tf_model/model_seq_1.keras', save_best_only=True)

model.compile(optimizer=Adam(learning_rate=0.0001), loss=BinaryCrossentropy(),
metrics=['Accuracy', AUC(name='auc')])

model.fit(X_train_pad, y_train, validation_data=(X_val_pad, y_val), epochs=20,
callbacks=[cp])
```

Epoch 1/20

167/167 5s 22ms/step -

Accuracy: 0.5982 - auc: 0.6246 - loss: 0.6557 - val\_Accuracy: 0.7609 - val\_auc:  
0.8161 - val\_loss: 0.5218

Epoch 2/20

167/167 4s 21ms/step -

Accuracy: 0.7674 - auc: 0.8276 - loss: 0.5050 - val\_Accuracy: 0.7763 - val\_auc:  
0.8482 - val\_loss: 0.4837

Epoch 3/20

167/167 4s 22ms/step -

Accuracy: 0.7987 - auc: 0.8520 - loss: 0.4593 - val\_Accuracy: 0.7973 - val\_auc:  
0.8635 - val\_loss: 0.4514

Epoch 4/20

167/167 4s 22ms/step -

Accuracy: 0.8052 - auc: 0.8683 - loss: 0.4383 - val\_Accuracy: 0.8074 - val\_auc:  
0.8687 - val\_loss: 0.4409

Epoch 5/20

167/167 4s 21ms/step -

Accuracy: 0.8015 - auc: 0.8634 - loss: 0.4428 - val\_Accuracy: 0.8008 - val\_auc:  
0.8709 - val\_loss: 0.4419

Epoch 6/20

167/167                    4s 21ms/step -  
 Accuracy: 0.8143 - auc: 0.8781 - loss: 0.4189 - val\_Accuracy: 0.8043 - val\_auc:  
 0.8721 - val\_loss: 0.4352  
 Epoch 7/20

167/167                    4s 21ms/step -  
 Accuracy: 0.8098 - auc: 0.8758 - loss: 0.4245 - val\_Accuracy: 0.8065 - val\_auc:  
 0.8736 - val\_loss: 0.4342  
 Epoch 8/20

167/167                    4s 21ms/step -  
 Accuracy: 0.8197 - auc: 0.8751 - loss: 0.4191 - val\_Accuracy: 0.7999 - val\_auc:  
 0.8743 - val\_loss: 0.4394  
 Epoch 9/20

167/167                    4s 22ms/step -  
 Accuracy: 0.8138 - auc: 0.8847 - loss: 0.4145 - val\_Accuracy: 0.8060 - val\_auc:  
 0.8738 - val\_loss: 0.4355  
 Epoch 10/20

167/167                    4s 21ms/step -  
 Accuracy: 0.8259 - auc: 0.8876 - loss: 0.4019 - val\_Accuracy: 0.8074 - val\_auc:  
 0.8743 - val\_loss: 0.4295  
 Epoch 11/20

167/167                    3s 21ms/step -  
 Accuracy: 0.8297 - auc: 0.8910 - loss: 0.3983 - val\_Accuracy: 0.8043 - val\_auc:  
 0.8742 - val\_loss: 0.4348  
 Epoch 12/20

167/167                    4s 21ms/step -  
 Accuracy: 0.8304 - auc: 0.8917 - loss: 0.3956 - val\_Accuracy: 0.8082 - val\_auc:  
 0.8750 - val\_loss: 0.4279  
 Epoch 13/20

167/167                    4s 22ms/step -  
 Accuracy: 0.8221 - auc: 0.8896 - loss: 0.3990 - val\_Accuracy: 0.8113 - val\_auc:  
 0.8764 - val\_loss: 0.4292  
 Epoch 14/20

167/167                    4s 21ms/step -  
 Accuracy: 0.8244 - auc: 0.8883 - loss: 0.4015 - val\_Accuracy: 0.8087 - val\_auc:  
 0.8755 - val\_loss: 0.4328  
 Epoch 15/20

167/167                    4s 22ms/step -  
 Accuracy: 0.8355 - auc: 0.8957 - loss: 0.3816 - val\_Accuracy: 0.8091 - val\_auc:  
 0.8756 - val\_loss: 0.4386  
 Epoch 16/20

167/167                    4s 22ms/step -  
 Accuracy: 0.8319 - auc: 0.9023 - loss: 0.3753 - val\_Accuracy: 0.8047 - val\_auc:  
 0.8752 - val\_loss: 0.4362  
 Epoch 17/20

167/167                    4s 22ms/step -  
 Accuracy: 0.8456 - auc: 0.9034 - loss: 0.3714 - val\_Accuracy: 0.8039 - val\_auc:  
 0.8741 - val\_loss: 0.4374  
 Epoch 18/20

```

167/167          4s 21ms/step -
Accuracy: 0.8465 - auc: 0.9062 - loss: 0.3674 - val_Accuracy: 0.8087 - val_auc:
0.8757 - val_loss: 0.4299
Epoch 19/20
167/167          4s 22ms/step -
Accuracy: 0.8403 - auc: 0.9040 - loss: 0.3722 - val_Accuracy: 0.8109 - val_auc:
0.8764 - val_loss: 0.4331
Epoch 20/20
167/167          4s 21ms/step -
Accuracy: 0.8406 - auc: 0.9005 - loss: 0.3760 - val_Accuracy: 0.8100 - val_auc:
0.8748 - val_loss: 0.4398

```

[19]: <keras.src.callbacks.history.History at 0x1781fe160>

Let us load the best model and run predictions on the test set

```

[100]: best_model = load_model('tf_model/model_seq_1.keras')


predictions = (best_model.predict(X_test_pad) > 0.5).astype(int)
predictions
output_df = pd.DataFrame(df_test['id'])
output_df['target'] = predictions
out = output_df.to_csv('Disaster_tweets_RNN_seq_1.csv', index=False)

```

```

102/102          1s 6ms/step

```

 Disaster\_tweets\_RNN\_seq\_1.csv  
Complete · 14s ago

The model got a score of 0.79773 which is not bad.

**Parameter tuning** Let us see if we can get a better score. We already know that the tweet ‘target’ - ie if a tweet is of a real disaster or not, is skewed with more zeros (not disaster) than ones (disaster). Let us use the weights parameter of the fit to see if we can get better results

```

[22]: disaster_counts = pd.Series(df_train['target']).value_counts()
disaster_counts
weights = {0:disaster_counts.sum()/disaster_counts[0], 1:disaster_counts.sum()/
↪disaster_counts[1]}
weights

```

[22]: {0: 1.753339474896361, 1: 2.3274228064811986}

```

[102]: model = Sequential([])

model.add(layers.Input(shape=(35,100)))
model.add(layers.LSTM(64, return_sequences=True))
model.add(layers.Dropout(0.2))
model.add(layers.LSTM(64, return_sequences=True))
model.add(layers.Dropout(0.2))
model.add(layers.Flatten())

```

```

model.add(layers.Dense(1, activation='sigmoid'))

cp = ModelCheckpoint('tf_model/model.keras', save_best_only=True)

model.compile(optimizer=Adam(learning_rate=0.0001), loss=BinaryCrossentropy(),
    metrics=['Accuracy', AUC(name='auc')])

model.fit(X_train_pad, y_train, validation_data=(X_val_pad, y_val), epochs=20,
    callbacks=[cp], class_weight=weights)

```

Epoch 1/20

167/167 5s 22ms/step -  
 Accuracy: 0.6667 - auc: 0.7015 - loss: 1.3042 - val\_Accuracy: 0.7487 - val\_auc:  
 0.8175 - val\_loss: 0.5237

Epoch 2/20

167/167 4s 22ms/step -  
 Accuracy: 0.7798 - auc: 0.8384 - loss: 0.9874 - val\_Accuracy: 0.7785 - val\_auc:  
 0.8495 - val\_loss: 0.4775

Epoch 3/20

167/167 4s 22ms/step -  
 Accuracy: 0.7965 - auc: 0.8611 - loss: 0.9091 - val\_Accuracy: 0.7885 - val\_auc:  
 0.8631 - val\_loss: 0.4644

Epoch 4/20

167/167 4s 22ms/step -  
 Accuracy: 0.7958 - auc: 0.8635 - loss: 0.9012 - val\_Accuracy: 0.7973 - val\_auc:  
 0.8647 - val\_loss: 0.4458

Epoch 5/20

167/167 4s 22ms/step -  
 Accuracy: 0.8015 - auc: 0.8702 - loss: 0.8887 - val\_Accuracy: 0.7999 - val\_auc:  
 0.8683 - val\_loss: 0.4511

Epoch 6/20

167/167 4s 23ms/step -  
 Accuracy: 0.8151 - auc: 0.8851 - loss: 0.8364 - val\_Accuracy: 0.7964 - val\_auc:  
 0.8695 - val\_loss: 0.4529

Epoch 7/20

167/167 4s 23ms/step -  
 Accuracy: 0.8107 - auc: 0.8763 - loss: 0.8670 - val\_Accuracy: 0.8060 - val\_auc:  
 0.8709 - val\_loss: 0.4363

Epoch 8/20

167/167 4s 23ms/step -  
 Accuracy: 0.8331 - auc: 0.8914 - loss: 0.8016 - val\_Accuracy: 0.8087 - val\_auc:  
 0.8718 - val\_loss: 0.4384

Epoch 9/20

167/167 4s 22ms/step -  
 Accuracy: 0.8189 - auc: 0.8824 - loss: 0.8407 - val\_Accuracy: 0.8082 - val\_auc:  
 0.8723 - val\_loss: 0.4376

Epoch 10/20

167/167 4s 23ms/step -



```

Accuracy: 0.8255 - auc: 0.8864 - loss: 0.8230 - val_Accuracy: 0.8065 - val_auc:
0.8729 - val_loss: 0.4346
Epoch 11/20
167/167          4s 24ms/step -
Accuracy: 0.8353 - auc: 0.8964 - loss: 0.7893 - val_Accuracy: 0.8060 - val_auc:
0.8733 - val_loss: 0.4388
Epoch 12/20
167/167          4s 23ms/step -
Accuracy: 0.8190 - auc: 0.8870 - loss: 0.8251 - val_Accuracy: 0.8117 - val_auc:
0.8732 - val_loss: 0.4329
Epoch 13/20
167/167          4s 23ms/step -
Accuracy: 0.8295 - auc: 0.8963 - loss: 0.7905 - val_Accuracy: 0.7968 - val_auc:
0.8732 - val_loss: 0.4462
Epoch 14/20
167/167          4s 23ms/step -
Accuracy: 0.8211 - auc: 0.8892 - loss: 0.8159 - val_Accuracy: 0.8100 - val_auc:
0.8733 - val_loss: 0.4352
Epoch 15/20
167/167          4s 23ms/step -
Accuracy: 0.8250 - auc: 0.8990 - loss: 0.7865 - val_Accuracy: 0.8065 - val_auc:
0.8735 - val_loss: 0.4326
Epoch 16/20
167/167          4s 23ms/step -
Accuracy: 0.8238 - auc: 0.8927 - loss: 0.7987 - val_Accuracy: 0.8052 - val_auc:
0.8734 - val_loss: 0.4330
Epoch 17/20
167/167          4s 23ms/step -
Accuracy: 0.8371 - auc: 0.9002 - loss: 0.7726 - val_Accuracy: 0.8065 - val_auc:
0.8738 - val_loss: 0.4317
Epoch 18/20
167/167          4s 23ms/step -
Accuracy: 0.8341 - auc: 0.8987 - loss: 0.7873 - val_Accuracy: 0.8087 - val_auc:
0.8742 - val_loss: 0.4338
Epoch 19/20
167/167          4s 22ms/step -
Accuracy: 0.8369 - auc: 0.9052 - loss: 0.7611 - val_Accuracy: 0.8078 - val_auc:
0.8738 - val_loss: 0.4383
Epoch 20/20
167/167          4s 24ms/step -
Accuracy: 0.8482 - auc: 0.9096 - loss: 0.7418 - val_Accuracy: 0.8039 - val_auc:
0.8706 - val_loss: 0.4405

```

```
[102]: <keras.src.callbacks.history.History at 0x3173fc0a0>
```

```
[103]: best_model = load_model('tf_model/model.keras')
```

```

predictions = (best_model.predict(X_test_pad) > 0.5).astype(int)
predictions
output_df = pd.DataFrame(df_test['id'])
output_df['target'] = predictions
out = output_df.to_csv('Disaster_tweets_RNN.csv', index=False)

```

102/102                      1s 6ms/step

 Disaster\_tweets\_RNN.csv  
Complete · 16s ago

0.80416

The model got a score of 0.80416

```

[104]: model = Sequential([])

model.add(layers.Input(shape=(35,100)))
model.add(layers.LSTM(128, return_sequences=True))
model.add(layers.Dropout(0.2))
model.add(layers.LSTM(128, return_sequences=True))
model.add(layers.Dropout(0.2))
model.add(layers.Flatten())
model.add(layers.Dense(1, activation='sigmoid'))

cp = ModelCheckpoint('tf_model/model_128.keras', save_best_only=True)

model.compile(optimizer=Adam(learning_rate=0.0001), loss=BinaryCrossentropy(),
↳ metrics=['Accuracy', AUC(name='auc')])

model.fit(X_train_pad, y_train, validation_data=(X_val_pad, y_val), epochs=50,
↳ callbacks=[cp], class_weight=weights)

best_model = load_model('tf_model/model_128.keras')

predictions = (best_model.predict(X_test_pad) > 0.5).astype(int)
predictions
output_df = pd.DataFrame(df_test['id'])
output_df['target'] = predictions
out = output_df.to_csv('Disaster_tweets_RNN_128.csv', index=False)

```

Epoch 1/50

167/167                      10s 53ms/step -

Accuracy: 0.7043 - auc: 0.7413 - loss: 1.2223 - val\_Accuracy: 0.7662 - val\_auc:  
0.8412 - val\_loss: 0.4967

Epoch 2/50

167/167                      10s 57ms/step -

Accuracy: 0.7967 - auc: 0.8566 - loss: 0.9263 - val\_Accuracy: 0.7863 - val\_auc:  
0.8581 - val\_loss: 0.4636

Epoch 3/50

167/167                      10s 59ms/step -

Accuracy: 0.8006 - auc: 0.8637 - loss: 0.9001 - val\_Accuracy: 0.7938 - val\_auc:

0.8662 - val\_loss: 0.4582  
 Epoch 4/50  
 167/167                    10s 61ms/step -  
 Accuracy: 0.8125 - auc: 0.8760 - loss: 0.8635 - val\_Accuracy: 0.8012 - val\_auc:  
 0.8689 - val\_loss: 0.4436  
 Epoch 5/50  
 167/167                    10s 60ms/step -  
 Accuracy: 0.8211 - auc: 0.8859 - loss: 0.8327 - val\_Accuracy: 0.8078 - val\_auc:  
 0.8704 - val\_loss: 0.4360  
 Epoch 6/50  
 167/167                    8s 48ms/step -  
 Accuracy: 0.8167 - auc: 0.8854 - loss: 0.8302 - val\_Accuracy: 0.8039 - val\_auc:  
 0.8713 - val\_loss: 0.4360  
 Epoch 7/50  
 167/167                    8s 45ms/step -  
 Accuracy: 0.8097 - auc: 0.8774 - loss: 0.8599 - val\_Accuracy: 0.8078 - val\_auc:  
 0.8716 - val\_loss: 0.4368  
 Epoch 8/50  
 167/167                    8s 46ms/step -  
 Accuracy: 0.8120 - auc: 0.8809 - loss: 0.8489 - val\_Accuracy: 0.8060 - val\_auc:  
 0.8725 - val\_loss: 0.4321  
 Epoch 9/50  
 167/167                    8s 45ms/step -  
 Accuracy: 0.8274 - auc: 0.8978 - loss: 0.7882 - val\_Accuracy: 0.8082 - val\_auc:  
 0.8737 - val\_loss: 0.4296  
 Epoch 10/50  
 167/167                    7s 45ms/step -  
 Accuracy: 0.8253 - auc: 0.8935 - loss: 0.8098 - val\_Accuracy: 0.8100 - val\_auc:  
 0.8725 - val\_loss: 0.4343  
 Epoch 11/50  
 167/167                    8s 45ms/step -  
 Accuracy: 0.8265 - auc: 0.8918 - loss: 0.7994 - val\_Accuracy: 0.8052 - val\_auc:  
 0.8707 - val\_loss: 0.4372  
 Epoch 12/50  
 167/167                    7s 43ms/step -  
 Accuracy: 0.8293 - auc: 0.9039 - loss: 0.7630 - val\_Accuracy: 0.8047 - val\_auc:  
 0.8732 - val\_loss: 0.4361  
 Epoch 13/50  
 167/167                    8s 47ms/step -  
 Accuracy: 0.8367 - auc: 0.9059 - loss: 0.7608 - val\_Accuracy: 0.8034 - val\_auc:  
 0.8740 - val\_loss: 0.4454  
 Epoch 14/50  
 167/167                    8s 48ms/step -  
 Accuracy: 0.8363 - auc: 0.9072 - loss: 0.7464 - val\_Accuracy: 0.8030 - val\_auc:  
 0.8712 - val\_loss: 0.4337  
 Epoch 15/50  
 167/167                    8s 47ms/step -  
 Accuracy: 0.8443 - auc: 0.9086 - loss: 0.7425 - val\_Accuracy: 0.8069 - val\_auc:

0.8721 - val\_loss: 0.4392  
 Epoch 16/50  
 167/167                8s 49ms/step -  
 Accuracy: 0.8388 - auc: 0.9080 - loss: 0.7453 - val\_Accuracy: 0.8082 - val\_auc:  
 0.8701 - val\_loss: 0.4461  
 Epoch 17/50  
 167/167                8s 48ms/step -  
 Accuracy: 0.8499 - auc: 0.9181 - loss: 0.7088 - val\_Accuracy: 0.7758 - val\_auc:  
 0.8667 - val\_loss: 0.5004  
 Epoch 18/50  
 167/167                8s 47ms/step -  
 Accuracy: 0.8458 - auc: 0.9195 - loss: 0.7044 - val\_Accuracy: 0.7995 - val\_auc:  
 0.8705 - val\_loss: 0.4676  
 Epoch 19/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.8538 - auc: 0.9239 - loss: 0.6787 - val\_Accuracy: 0.8082 - val\_auc:  
 0.8674 - val\_loss: 0.4491  
 Epoch 20/50  
 167/167                8s 47ms/step -  
 Accuracy: 0.8633 - auc: 0.9296 - loss: 0.6547 - val\_Accuracy: 0.7929 - val\_auc:  
 0.8668 - val\_loss: 0.4982  
 Epoch 21/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.8676 - auc: 0.9341 - loss: 0.6358 - val\_Accuracy: 0.8052 - val\_auc:  
 0.8652 - val\_loss: 0.4538  
 Epoch 22/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.8629 - auc: 0.9317 - loss: 0.6540 - val\_Accuracy: 0.8060 - val\_auc:  
 0.8628 - val\_loss: 0.4700  
 Epoch 23/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.8727 - auc: 0.9389 - loss: 0.6127 - val\_Accuracy: 0.7973 - val\_auc:  
 0.8611 - val\_loss: 0.5080  
 Epoch 24/50  
 167/167                8s 47ms/step -  
 Accuracy: 0.8736 - auc: 0.9429 - loss: 0.5980 - val\_Accuracy: 0.8043 - val\_auc:  
 0.8585 - val\_loss: 0.4746  
 Epoch 25/50  
 167/167                8s 48ms/step -  
 Accuracy: 0.8803 - auc: 0.9515 - loss: 0.5528 - val\_Accuracy: 0.7968 - val\_auc:  
 0.8550 - val\_loss: 0.5123  
 Epoch 26/50  
 167/167                8s 47ms/step -  
 Accuracy: 0.8896 - auc: 0.9557 - loss: 0.5266 - val\_Accuracy: 0.7771 - val\_auc:  
 0.8533 - val\_loss: 0.5466  
 Epoch 27/50  
 167/167                8s 45ms/step -  
 Accuracy: 0.8983 - auc: 0.9563 - loss: 0.5196 - val\_Accuracy: 0.7842 - val\_auc:

0.8515 - val\_loss: 0.5469  
 Epoch 28/50  
 167/167                8s 48ms/step -  
 Accuracy: 0.8944 - auc: 0.9614 - loss: 0.4960 - val\_Accuracy: 0.7947 - val\_auc:  
 0.8453 - val\_loss: 0.5197  
 Epoch 29/50  
 167/167                8s 47ms/step -  
 Accuracy: 0.8933 - auc: 0.9619 - loss: 0.4933 - val\_Accuracy: 0.7925 - val\_auc:  
 0.8504 - val\_loss: 0.5921  
 Epoch 30/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.9062 - auc: 0.9662 - loss: 0.4565 - val\_Accuracy: 0.7977 - val\_auc:  
 0.8488 - val\_loss: 0.5655  
 Epoch 31/50  
 167/167                8s 47ms/step -  
 Accuracy: 0.9024 - auc: 0.9633 - loss: 0.4761 - val\_Accuracy: 0.7793 - val\_auc:  
 0.8493 - val\_loss: 0.6162  
 Epoch 32/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.9111 - auc: 0.9707 - loss: 0.4249 - val\_Accuracy: 0.7929 - val\_auc:  
 0.8390 - val\_loss: 0.6489  
 Epoch 33/50  
 167/167                8s 47ms/step -  
 Accuracy: 0.9128 - auc: 0.9725 - loss: 0.4147 - val\_Accuracy: 0.7986 - val\_auc:  
 0.8456 - val\_loss: 0.6494  
 Epoch 34/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.9267 - auc: 0.9788 - loss: 0.3686 - val\_Accuracy: 0.7868 - val\_auc:  
 0.8402 - val\_loss: 0.7254  
 Epoch 35/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.9224 - auc: 0.9752 - loss: 0.3915 - val\_Accuracy: 0.7824 - val\_auc:  
 0.8404 - val\_loss: 0.6878  
 Epoch 36/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.9308 - auc: 0.9818 - loss: 0.3408 - val\_Accuracy: 0.7820 - val\_auc:  
 0.8381 - val\_loss: 0.7623  
 Epoch 37/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.9274 - auc: 0.9821 - loss: 0.3325 - val\_Accuracy: 0.7789 - val\_auc:  
 0.8414 - val\_loss: 0.7350  
 Epoch 38/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.9313 - auc: 0.9817 - loss: 0.3361 - val\_Accuracy: 0.7434 - val\_auc:  
 0.8343 - val\_loss: 0.7677  
 Epoch 39/50  
 167/167                8s 46ms/step -  
 Accuracy: 0.9313 - auc: 0.9834 - loss: 0.3292 - val\_Accuracy: 0.7785 - val\_auc:

```

0.8368 - val_loss: 0.7608
Epoch 40/50
167/167      8s 46ms/step -
Accuracy: 0.9378 - auc: 0.9834 - loss: 0.3243 - val_Accuracy: 0.7557 - val_auc:
0.8331 - val_loss: 0.8572
Epoch 41/50
167/167      8s 46ms/step -
Accuracy: 0.9417 - auc: 0.9885 - loss: 0.2727 - val_Accuracy: 0.7789 - val_auc:
0.8294 - val_loss: 0.8503
Epoch 42/50
167/167      8s 46ms/step -
Accuracy: 0.9475 - auc: 0.9892 - loss: 0.2613 - val_Accuracy: 0.7780 - val_auc:
0.8210 - val_loss: 0.8456
Epoch 43/50
167/167      8s 46ms/step -
Accuracy: 0.9474 - auc: 0.9882 - loss: 0.2720 - val_Accuracy: 0.7837 - val_auc:
0.8336 - val_loss: 0.8340
Epoch 44/50
167/167      8s 46ms/step -
Accuracy: 0.9407 - auc: 0.9850 - loss: 0.3032 - val_Accuracy: 0.7741 - val_auc:
0.8287 - val_loss: 0.9206
Epoch 45/50
167/167      8s 46ms/step -
Accuracy: 0.9542 - auc: 0.9919 - loss: 0.2310 - val_Accuracy: 0.7732 - val_auc:
0.8278 - val_loss: 0.8969
Epoch 46/50
167/167      8s 46ms/step -
Accuracy: 0.9535 - auc: 0.9901 - loss: 0.2515 - val_Accuracy: 0.7509 - val_auc:
0.8221 - val_loss: 0.9952
Epoch 47/50
167/167      8s 48ms/step -
Accuracy: 0.9437 - auc: 0.9880 - loss: 0.2827 - val_Accuracy: 0.7688 - val_auc:
0.8240 - val_loss: 0.9393
Epoch 48/50
167/167      8s 45ms/step -
Accuracy: 0.9599 - auc: 0.9928 - loss: 0.2146 - val_Accuracy: 0.7912 - val_auc:
0.8226 - val_loss: 0.9998
Epoch 49/50
167/167      8s 47ms/step -
Accuracy: 0.9612 - auc: 0.9941 - loss: 0.1952 - val_Accuracy: 0.7815 - val_auc:
0.8297 - val_loss: 0.9987
Epoch 50/50
167/167      8s 46ms/step -
Accuracy: 0.9569 - auc: 0.9931 - loss: 0.2100 - val_Accuracy: 0.7780 - val_auc:
0.8234 - val_loss: 0.9592
102/102      1s 12ms/step
This model with 128 units and 50 epochs doesnt perform

```

any better than the model with 64 units and 20 epochs.



Disaster\_tweets\_RNN\_128.csv  
Complete · 18s ago

0.80140

Changing the learning rate to 0.00001 (not captured here) also didn't help with increasing the accuracy

Let us replace LSTM with GRU

```
[108]: model = Sequential([])
model.add(layers.Input(shape=(35,100)))
model.add(layers.GRU(64, return_sequences=True))
model.add(layers.Dropout(0.2))
model.add(layers.GRU(64, return_sequences=True))
model.add(layers.Dropout(0.2))
model.add(layers.Flatten())
model.add(layers.Dense(1, activation='sigmoid'))

cp = ModelCheckpoint('tf_model/model_gru.keras', save_best_only=True)

model.compile(optimizer=Adam(learning_rate=0.01), loss=BinaryCrossentropy(),
↳metrics=['Accuracy'])

model.fit(X_train_pad, y_train, validation_data=(X_val_pad, y_val), epochs=30,
↳callbacks=[cp], class_weight=weights, batch_size=512)

best_model = load_model('tf_model/model_gru.keras')

predictions = (best_model.predict(X_test_pad) > 0.5).astype(int)
predictions
output_df = pd.DataFrame(df_test['id'])
output_df['target'] = predictions
out = output_df.to_csv('Disaster_tweets_RNN_gru.csv', index=False)
```

Epoch 1/30

11/11 3s 191ms/step -

Accuracy: 0.5691 - loss: 1.4738 - val\_Accuracy: 0.7680 - val\_loss: 0.5020

Epoch 2/30

11/11 2s 185ms/step -

Accuracy: 0.7757 - loss: 0.9967 - val\_Accuracy: 0.7802 - val\_loss: 0.4829

Epoch 3/30

11/11 2s 191ms/step -

Accuracy: 0.8056 - loss: 0.8889 - val\_Accuracy: 0.7881 - val\_loss: 0.4583

Epoch 4/30

11/11 2s 190ms/step -

Accuracy: 0.8186 - loss: 0.8311 - val\_Accuracy: 0.8069 - val\_loss: 0.4358

Epoch 5/30

11/11 2s 194ms/step -

Accuracy: 0.8335 - loss: 0.7582 - val\_Accuracy: 0.8122 - val\_loss: 0.4369

Epoch 6/30  
 11/11 2s 197ms/step -  
 Accuracy: 0.8640 - loss: 0.6833 - val\_Accuracy: 0.7846 - val\_loss: 0.4692

Epoch 7/30  
 11/11 2s 193ms/step -  
 Accuracy: 0.8720 - loss: 0.6241 - val\_Accuracy: 0.7947 - val\_loss: 0.5146

Epoch 8/30  
 11/11 2s 198ms/step -  
 Accuracy: 0.8821 - loss: 0.5609 - val\_Accuracy: 0.8012 - val\_loss: 0.5165

Epoch 9/30  
 11/11 2s 200ms/step -  
 Accuracy: 0.8949 - loss: 0.4900 - val\_Accuracy: 0.7693 - val\_loss: 0.6139

Epoch 10/30  
 11/11 2s 188ms/step -  
 Accuracy: 0.9224 - loss: 0.3883 - val\_Accuracy: 0.7806 - val\_loss: 0.7008

Epoch 11/30  
 11/11 2s 202ms/step -  
 Accuracy: 0.9487 - loss: 0.2721 - val\_Accuracy: 0.7758 - val\_loss: 0.8283

Epoch 12/30  
 11/11 2s 193ms/step -  
 Accuracy: 0.9579 - loss: 0.2330 - val\_Accuracy: 0.7793 - val\_loss: 0.8612

Epoch 13/30  
 11/11 2s 188ms/step -  
 Accuracy: 0.9465 - loss: 0.2846 - val\_Accuracy: 0.7614 - val\_loss: 0.7830

Epoch 14/30  
 11/11 2s 185ms/step -  
 Accuracy: 0.9607 - loss: 0.2194 - val\_Accuracy: 0.7680 - val\_loss: 0.9275

Epoch 15/30  
 11/11 2s 191ms/step -  
 Accuracy: 0.9682 - loss: 0.1822 - val\_Accuracy: 0.7785 - val\_loss: 0.8825

Epoch 16/30  
 11/11 2s 197ms/step -  
 Accuracy: 0.9780 - loss: 0.1417 - val\_Accuracy: 0.7785 - val\_loss: 1.0798

Epoch 17/30  
 11/11 2s 204ms/step -  
 Accuracy: 0.9803 - loss: 0.1145 - val\_Accuracy: 0.7763 - val\_loss: 1.0643

Epoch 18/30  
 11/11 2s 194ms/step -  
 Accuracy: 0.9802 - loss: 0.1304 - val\_Accuracy: 0.7837 - val\_loss: 1.0214

Epoch 19/30  
 11/11 2s 196ms/step -  
 Accuracy: 0.9779 - loss: 0.1360 - val\_Accuracy: 0.7780 - val\_loss: 1.0073

Epoch 20/30  
 11/11 2s 192ms/step -  
 Accuracy: 0.9854 - loss: 0.1095 - val\_Accuracy: 0.7771 - val\_loss: 1.0967

Epoch 21/30  
 11/11 2s 186ms/step -  
 Accuracy: 0.9786 - loss: 0.1163 - val\_Accuracy: 0.7758 - val\_loss: 1.1032



```

Epoch 22/30
11/11          2s 192ms/step -
Accuracy: 0.9839 - loss: 0.0935 - val_Accuracy: 0.7697 - val_loss: 1.0899
Epoch 23/30
11/11          2s 187ms/step -
Accuracy: 0.9827 - loss: 0.0883 - val_Accuracy: 0.7644 - val_loss: 1.0857
Epoch 24/30
11/11          2s 197ms/step -
Accuracy: 0.9831 - loss: 0.0804 - val_Accuracy: 0.7701 - val_loss: 1.2210
Epoch 25/30
11/11          2s 196ms/step -
Accuracy: 0.9833 - loss: 0.0944 - val_Accuracy: 0.7583 - val_loss: 1.0229
Epoch 26/30
11/11          2s 199ms/step -
Accuracy: 0.9845 - loss: 0.0820 - val_Accuracy: 0.7758 - val_loss: 1.2088
Epoch 27/30
11/11          2s 193ms/step -
Accuracy: 0.9792 - loss: 0.1074 - val_Accuracy: 0.7815 - val_loss: 1.0998
Epoch 28/30
11/11          2s 202ms/step -
Accuracy: 0.9852 - loss: 0.0779 - val_Accuracy: 0.7732 - val_loss: 1.2302
Epoch 29/30
11/11          2s 195ms/step -
Accuracy: 0.9834 - loss: 0.0764 - val_Accuracy: 0.7715 - val_loss: 1.2893
Epoch 30/30
11/11          2s 193ms/step -
Accuracy: 0.9859 - loss: 0.0661 - val_Accuracy: 0.7728 - val_loss: 1.3217
102/102        1s 5ms/step

```

With learning rate reduced to 0.01, replacing LSTM with GRU and validation to just Accuracy, batch size at 512 and epochs at 30, the model is slightly better at 0.81152



0.81152

## 0.5 Step 4: Results and Analysis

```

[23]: from tensorflow.keras.callbacks import TensorBoard

model = Sequential([])
model.add(layers.Input(shape=(35,100)))
model.add(layers.GRU(64, return_sequences=True))
model.add(layers.Dropout(0.2))
model.add(layers.GRU(64, return_sequences=True))
model.add(layers.Dropout(0.2))
model.add(layers.Flatten())
model.add(layers.Dense(1, activation='sigmoid'))

log_dir = 'tf_model_gru'

```

```

tensorboard_callback = TensorBoard(log_dir=log_dir, histogram_freq=1)

model.compile(optimizer=Adam(learning_rate=0.01), loss=BinaryCrossentropy(),  

↳metrics=['Accuracy'])

model.fit(X_train_pad, y_train, validation_data=(X_val_pad, y_val), epochs=30,  

↳callbacks=[tensorboard_callback], class_weight=weights, batch_size=512)

```

Epoch 1/30

11/11 3s 181ms/step -

Accuracy: 0.5639 - loss: 1.5175 - val\_Accuracy: 0.7566 - val\_loss: 0.5090

Epoch 2/30

11/11 2s 180ms/step -

Accuracy: 0.7842 - loss: 0.9594 - val\_Accuracy: 0.7833 - val\_loss: 0.4672

Epoch 3/30

11/11 2s 189ms/step -

Accuracy: 0.7933 - loss: 0.9042 - val\_Accuracy: 0.8104 - val\_loss: 0.4425

Epoch 4/30

11/11 2s 187ms/step -

Accuracy: 0.8169 - loss: 0.8274 - val\_Accuracy: 0.8122 - val\_loss: 0.4328

Epoch 5/30

11/11 2s 190ms/step -

Accuracy: 0.8281 - loss: 0.7867 - val\_Accuracy: 0.7885 - val\_loss: 0.4702

Epoch 6/30

11/11 2s 187ms/step -

Accuracy: 0.8466 - loss: 0.7219 - val\_Accuracy: 0.7890 - val\_loss: 0.4732

Epoch 7/30

11/11 2s 187ms/step -

Accuracy: 0.8520 - loss: 0.6713 - val\_Accuracy: 0.7999 - val\_loss: 0.4557

Epoch 8/30

11/11 2s 193ms/step -

Accuracy: 0.8825 - loss: 0.6105 - val\_Accuracy: 0.7920 - val\_loss: 0.4821

Epoch 9/30

11/11 2s 190ms/step -

Accuracy: 0.9089 - loss: 0.4950 - val\_Accuracy: 0.7824 - val\_loss: 0.6624

Epoch 10/30

11/11 2s 194ms/step -

Accuracy: 0.9150 - loss: 0.4480 - val\_Accuracy: 0.7868 - val\_loss: 0.6272

Epoch 11/30

11/11 2s 191ms/step -

Accuracy: 0.9420 - loss: 0.3299 - val\_Accuracy: 0.7890 - val\_loss: 0.6917

Epoch 12/30

11/11 2s 202ms/step -

Accuracy: 0.9559 - loss: 0.2587 - val\_Accuracy: 0.7601 - val\_loss: 0.8525

Epoch 13/30

11/11 2s 195ms/step -

Accuracy: 0.9517 - loss: 0.2490 - val\_Accuracy: 0.7771 - val\_loss: 0.7325



Epoch 14/30

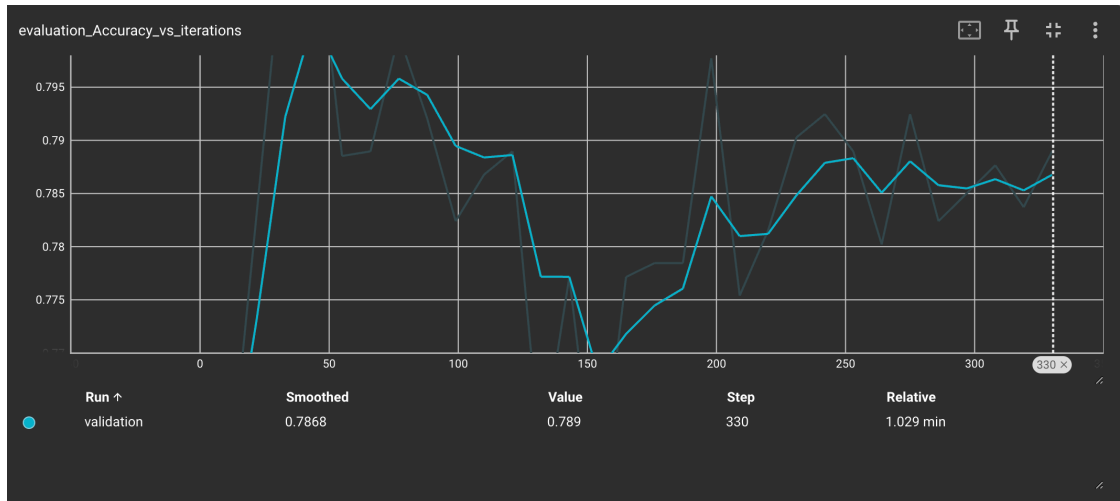
11/11                    2s 195ms/step -  
 Accuracy: 0.9666 - loss: 0.1905 - val\_Accuracy: 0.7548 - val\_loss: 0.8835  
 Epoch 15/30  
 11/11                    2s 191ms/step -  
 Accuracy: 0.9602 - loss: 0.2172 - val\_Accuracy: 0.7771 - val\_loss: 0.8118  
 Epoch 16/30  
 11/11                    2s 189ms/step -  
 Accuracy: 0.9693 - loss: 0.1879 - val\_Accuracy: 0.7785 - val\_loss: 0.9077  
 Epoch 17/30  
 11/11                    2s 192ms/step -  
 Accuracy: 0.9811 - loss: 0.1292 - val\_Accuracy: 0.7785 - val\_loss: 0.9023  
 Epoch 18/30  
 11/11                    2s 195ms/step -  
 Accuracy: 0.9790 - loss: 0.1280 - val\_Accuracy: 0.7977 - val\_loss: 1.0260  
 Epoch 19/30  
 11/11                    2s 195ms/step -  
 Accuracy: 0.9816 - loss: 0.1095 - val\_Accuracy: 0.7754 - val\_loss: 0.9692  
 Epoch 20/30  
 11/11                    2s 202ms/step -  
 Accuracy: 0.9801 - loss: 0.1165 - val\_Accuracy: 0.7815 - val\_loss: 1.0063  
 Epoch 21/30  
 11/11                    2s 193ms/step -  
 Accuracy: 0.9822 - loss: 0.0910 - val\_Accuracy: 0.7903 - val\_loss: 1.0736  
 Epoch 22/30  
 11/11                    2s 196ms/step -  
 Accuracy: 0.9813 - loss: 0.1086 - val\_Accuracy: 0.7925 - val\_loss: 1.0481  
 Epoch 23/30  
 11/11                    2s 206ms/step -  
 Accuracy: 0.9819 - loss: 0.1089 - val\_Accuracy: 0.7890 - val\_loss: 0.9706  
 Epoch 24/30  
 11/11                    2s 206ms/step -  
 Accuracy: 0.9800 - loss: 0.1036 - val\_Accuracy: 0.7802 - val\_loss: 1.0200  
 Epoch 25/30  
 11/11                    2s 209ms/step -  
 Accuracy: 0.9810 - loss: 0.0941 - val\_Accuracy: 0.7925 - val\_loss: 1.1031  
 Epoch 26/30  
 11/11                    2s 197ms/step -  
 Accuracy: 0.9823 - loss: 0.0886 - val\_Accuracy: 0.7824 - val\_loss: 1.1613  
 Epoch 27/30  
 11/11                    2s 190ms/step -  
 Accuracy: 0.9759 - loss: 0.0970 - val\_Accuracy: 0.7850 - val\_loss: 1.0427  
 Epoch 28/30  
 11/11                    2s 198ms/step -  
 Accuracy: 0.9862 - loss: 0.0787 - val\_Accuracy: 0.7877 - val\_loss: 1.1444  
 Epoch 29/30  
 11/11                    2s 211ms/step -  
 Accuracy: 0.9861 - loss: 0.0759 - val\_Accuracy: 0.7837 - val\_loss: 1.1274  
 Epoch 30/30

11/11                      2s 206ms/step -  
Accuracy: 0.9809 - loss: 0.0842 - val\_Accuracy: 0.7890 - val\_loss: 1.1987

[23]: <keras.src.callbacks.history.History at 0x3135b8c40>

TensorBoard for the GRU model:





We have explored LSTM and GRU based RNN. For the models proposed above, the models have returned results between 0.79-0.81. Due to lack of time, and since the primary expectation from this project is not accuracy, I am settling for these values.

A brief overview of what was attempted: | RNN family | units | learning rate | epochs | weights | batch size | metrics | test result | | :—: | :—: | :—: | :—: | :—: | :—: | :—: | :—: | :—: | :—: | | LSTM | 64 | 0.0001 | 20 | No | default | ['Accuracy', AUC(name='auc')] | 0.79773 | | LSTM | 64 | 0.0001 | 20 | Yes | default | ['Accuracy', AUC(name='auc')] | 0.80416 | | LSTM | 128 | 0.0001 | 50 | Yes | default | ['Accuracy', AUC(name='auc')] | 0.80140 | | GRU | 64 | 0.01 | 30 | Yes | 512 | ['Accuracy'] | 0.81152 |

metrics 'Accuracy' works well enough for most cases as has been seen in the results. I have explored fine tuning hyper-parameters and other values to see how the model reacts for both validation as well as test data. The values I used for learning rate or epoch does play a role in the accuracy but the results were not with a big margin.

I haven't shown all the models I tried since a few of the tweaks in hyper parameters didn't yield better results. Some of those models performed worse than the above models.

Decreasing the learning rate further to 0.00001 didn't help with the accuracy. epochs at 20 or 50 are also not really much helpful in these models. I could have probably gotten similar results with a lesser epoch.

## 0.6 Step 5: Conclusion

**Results and key learnings** The key learning is that tweaking hyperparameter alone is not enough in many cases. We need to fine tune the model based on the kind of data that is being used in the models. We could clean up the data a little more including removal of http links and explore further stop words removal. While the results are decent, unfortunately I would have been happy if I had come up with a model with a result  $> 0.9$

**Try in future** I have used only Adam optimizer and would like to try other optimizers to see how they perform for this use case. I have used GloVe embeddings and would try other word embeddings

## 0.7 Step 6: Produce Deliverables

GitHub Repository: [https://github.com/krishnakuruvadi/Week4\\_RNN](https://github.com/krishnakuruvadi/Week4_RNN)

Report: [https://github.com/krishnakuruvadi/Week4\\_RNN](https://github.com/krishnakuruvadi/Week4_RNN)

### Submissions

		Recent ▾
Submission and Description		Public Score ⓘ
✓	<b>Disaster_tweets_RNN_gru.csv</b> Complete · 21s ago	<b>0.81152</b>
✓	<b>Disaster_tweets_RNN_128.csv</b> Complete · 34m ago	<b>0.80140</b>
✓	<b>Disaster_tweets_RNN_seq_1.csv</b> Complete · 7h ago	<b>0.79773</b>
✓	<b>Disaster_tweets_RNN.csv</b> Complete · 21h ago	<b>0.80416</b>

Kaggle leaderboard:

[ ]: